

## Supervised Logistic Regression for Classification

---

### 0. Import library

```
1 # Import libraries
2
3 # math library
4 import numpy as np
5
6 # visualization library
7 %matplotlib inline
8 from IPython.display import set_matplotlib_formats
9 set_matplotlib_formats('png2x','pdf')
10 import matplotlib.pyplot as plt
11
12 # machine learning library
13 from sklearn.linear_model import LogisticRegression
14
15 # 3d visualization
16 from mpl_toolkits.mplot3d import axes3d
17
18 # computational time
19 import time
20
```

### 1. Load dataset

---

The data features  $x_i = (x_{i(1)}, x_{i(2)})$  represent 2 exam grades  $x_{i(1)}$  and  $x_{i(2)}$  for each student  $i$ .

The data label  $y_i$  indicates if the student  $i$  was admitted (value is 1) or rejected (value is 0).

```
1 # import data with numpy
2 data = np.loadtxt('/content/drive/My Drive/Colab Notebooks/MachineLearningProject/04/dataset.txt')
3
4 # number of training data
5 n = data.shape[0]
6 print('Number of training data=',n)
```

➤ Number of training data= 100

### 2. Explore the dataset distribution

---

Plot the training data points.


You may use matplotlib function `scatter(x,y)`.

```
1 x1 = data[:,0].astype(np.float128) # exam grade 1
2 x2 = data[:,1].astype(np.float128) # exam grade 2
3 idx = data[:,2].astype(np.float128)
4 #idx_admit = (idx==1) # index of students who were admitted
5 #idx_rejec = (idx==0) # index of students who were rejected
6
7 x1_idx0 = x1[idx == 0]
8 x1_idx1 = x1[idx == 1]
9
10 x2_idx0 = x2[idx == 0]
11 x2_idx1 = x2[idx == 1]
12
```

```

12
13 plt.figure(figsize=(8,8))
14 plt.scatter(x1_idx0, x2_idx0, alpha=0.5, c='b', label='Rejected')
15 plt.scatter(x1_idx1, x2_idx1, alpha=0.5, c='r', label='Admitted')
16 plt.title('Training data')
17 plt.xlabel('Exam grade 1')
18 plt.ylabel('Exam grade 2')
19 plt.legend()
20 plt.show()

```



### 3. Sigmoid/logistic function

$$\sigma(\eta) = \frac{1}{1 + \exp^{-\eta}}$$

Define and plot the sigmoid function for values in [-10,10]:


You may use functions `np.exp`, `np.linspace`.

```

1 def sigmoid(z):
2     try:
3         return 1 / (1 + np.exp(-z))
4     except OverflowError:
5         return 1e-9
6
7
8 # plot
9 x_values = np.linspace(-10,10)
10
11 plt.figure(2)
12 plt.plot(x_values,sigmoid(x_values))
13 plt.title("Sigmoid function")
14 plt.grid(True)

```





#### 4. Define the prediction function for the classification

The prediction function is defined by:

$$p_w(x) = \sigma(w_0 + w_1x_{(1)} + w_2x_{(2)}) = \sigma(w^T x)$$

Implement the prediction function in a vectorised way as follows:

$$X = \begin{bmatrix} 1 & x_{1(1)} & x_{1(2)} \\ 1 & x_{2(1)} & x_{2(2)} \\ \vdots & \vdots & \vdots \\ 1 & x_{n(1)} & x_{n(2)} \end{bmatrix} \quad \text{and} \quad w = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \end{bmatrix} \quad \Rightarrow \quad p_w(x) = \sigma(Xw) = \begin{bmatrix} \sigma(w_0 + w_1x_{1(1)} + w_2x_{1(2)}) \\ \sigma(w_0 + w_1x_{2(1)} + w_2x_{2(2)}) \\ \vdots \\ \sigma(w_0 + w_1x_{n(1)} + w_2x_{n(2)}) \end{bmatrix}$$

Use the new function `sigmoid`.

```
1 # construct the data matrix X
2 n = len(data)
3 X = np.array([[1,x1,x2] for x1,x2 in zip(x1, x2)])
4
5 # parameters vector
6 w = np.array([[1.0],[1.0],[1.0]])
7
8 # predictive function definition
9 def f_pred(X,w):
10
11     p = np.dot(X,w)
12
13     return p
14
15 y_pred = f_pred(X,w)
```

#### 5. Define the classification loss function

Mean Square Error

$$L(w) = \frac{1}{n} \sum_{i=1}^n (\sigma(w^T x_i) - y_i)^2$$

Cross-Entropy

$$L(w) = \frac{1}{n} \sum_{i=1}^n (-y_i \log(\sigma(w^T x_i)) - (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

The vectorized representation is for the mean square error is as follows:

$$L(w) = \frac{1}{n} (p_w(x) - y)^T (p_w(x) - y)$$

The vectorized representation is for the cross-entropy error is as follows:

$$L(w) = \frac{1}{n} \left( -y^T \log(p_w(x)) - (1 - y)^T \log(1 - p_w(x)) \right)$$

where

$$p_w(x) = \sigma(Xw) = \begin{bmatrix} \sigma(w_0 + w_1x_{1(1)} + w_2x_{1(2)}) \\ \sigma(w_0 + w_1x_{2(1)} + w_2x_{2(2)}) \\ \vdots \\ \sigma(w_0 + w_1x_{n(1)} + w_2x_{n(2)}) \end{bmatrix} \quad \text{and} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

You may use numpy functions `.T` and `np.log`.

```

1 def mse_loss(label, h_arr): # mean square error
2     loss = (np.dot((h_arr - label).T, (h_arr - label))) / len(h_arr)
3     return loss
4
5 def ce_loss(label, h_arr): # cross-entropy error
6     loss = -(np.dot(label.T, np.log(h_arr)) + np.dot((1-label).T, np.log(1-h_arr))) / len(h_arr)
7     return loss

```

## 6. Define the gradient of the classification loss function

Given the mean square loss

$$L(w) = \frac{1}{n} (p_w(x) - y)^T (p_w(x) - y)$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T ((p_w(x) - y) \odot (p_w(x) \odot (1 - p_w(x))))$$

Given the cross-entropy loss

$$L(w) = \frac{1}{n} (-y^T \log(p_w(x)) - (1 - y)^T \log(1 - p_w(x)))$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T (p_w(x) - y)$$

Implement the vectorized version of the gradient of the classification loss function

```

1 # loss function definition
2 def loss_logreg(y_pred,y):
3
4     n = len(y)
5     #loss = (np.dot((sigmoid(y_pred) - y).T, (sigmoid(y_pred) - y))) / n
6     loss = -(np.dot(y.T, np.log(sigmoid(y_pred))) + np.dot((1-y).T, np.log(1-sigmoid(y_pred)))) / n
7     return loss
8
9
10 # Test loss function
11 y = data[:,2][:,None] # label
12 y_pred = f_pred(X,w) # prediction
13
14 loss = loss_logreg(y_pred,y)

```

```

1 # gradient functions
2 def grad_loss(y_pred, y, X):
3     #grad = 2 * np.dot(X.T, np.dot((sigmoid(y_pred)-y), np.dot(sigmoid(y_pred).T, (1-sigmoid(y_pred)))) / n
4     grad = 2 * np.dot(X.T, (sigmoid(y_pred) - y)) / len(y_pred)
5     return grad

```

## 7. Implement the gradient descent algorithm

Vectorized implementation for the mean square loss:


$$w^{k+1} = w^k - \tau \frac{2}{n} X^T \left( (p_w(x) - y) \odot (p_w(x) \odot (1 - p_w(x))) \right)$$

Vectorized implementation for the cross-entropy loss:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T (p_w(x) - y)$$

Plot the loss values  $L(w^k)$  w.r.t. iteration  $k$  the number of iterations for the both loss functions.

```
1 # gradient descent function definition
2 def grad_desc(X, y , w_init=np.array([0,0,0])[:,None] ,tau=1e-4, max_iter=500):
3
4     L_iters = np.zeros([max_iter]) # record the loss values
5     w_iters = np.zeros([max_iter,2]) # record the loss values
6     w = w_init # initialization
7
8     for i in range(max_iter): # loop over the iterations
9
10        y_pred = f_pred(X,w) # linear predication function
11        grad_f = grad_loss(y_pred,y,X) # gradient of the loss
12        w = w - tau* grad_f # update rule of gradient descent
13        L_iters[i] = loss_logreg(y_pred,y) # save the current loss value
14        w_iters[i,:] = w[0],w[1] # save the current w value
15
16    return w, L_iters, w_iters
17
18
19 # run gradient descent algorithm
20 start = time.time()
21 w_init = np.array([-5,0,0])[:,None]
22 tau = 0.00000005; max_iter = 20000
23 w, L_iters, w_iters = grad_desc(X,y,w_init,tau,max_iter)
24
25
26 # plot
27 plt.figure(3)
28 plt.plot([op for op in range(max_iter)], L_iters, c='blue') # plot the loss curve
29 plt.xlabel('Iterations')
30 plt.ylabel('Loss value')
31 plt.show()
32
```



## 8. Plot the decision boundary

The decision boundary is defined by all points


$$x = (x_{(1)}, x_{(2)}) \quad \text{such that} \quad p_w(x) = 0.5$$

You may use numpy and matplotlib functions `np.meshgrid`, `np.linspace`, `reshape`, `contour`.

```

1 # compute values p(x) for multiple data points x
2 x1_min, x1_max = X[:,1].min(), X[:,1].max() # min and max of grade 1
3 x2_min, x2_max = X[:,2].min(), X[:,2].max() # min and max of grade 2
4 xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max)) # create meshgrid
5 X2 = np.ones([np.prod(xx1.shape),3])
6 X2[:,1] = xx1.reshape(-1)
7 X2[:,2] = xx2.reshape(-1)
8 p = f_pred(X2,w)
9 p = p.reshape(50,50)
10
11
12 # plot
13 plt.figure(4,figsize=(6,6))
14 plt.scatter(x1_idx0, x2_idx0, alpha=0.5, c='b', label='Rejected')
15 plt.scatter(x1_idx1, x2_idx1, alpha=0.5, c='r', label='Admitted')
16 plt.contour(xx1, xx2, p, 0)
17 plt.xlabel('Exam grade 1')
18 plt.ylabel('Exam grade 2')
19 plt.legend()
20 plt.title('Decision boundary')
21 plt.show()

```



## 9. Comparison with Scikit-learn logistic regression algorithm with the gradient descent with the cross-entropy loss

You may use scikit-learn function `LogisticRegression(C=1e6)`.

```

1 # run logistic regression with scikit-learn
2 start = time.time()
3 logreg_sklearn = LogisticRegression(C=1e6, random_state=3) # scikit-learn logistic regression
4 logreg_sklearn.fit(X, y) # learn the model parameters
5
6 # compute loss value
7 w_sklearn = np.zeros([3,1])
8 w_sklearn[0,0] = logreg_sklearn.predict(X)[0]
9 w_sklearn[1:3,0] = logreg_sklearn.predict(X)[1:3]
10 #loss_sklearn = loss_logreg( )
11
12 # plot

```

```


13 plt.figure(4,figsize=(6,6))
14 plt.scatter(x1_idx0, x2_idx0, alpha=0.5, c='b', label='Rejected')
15 plt.scatter(x1_idx1, x2_idx1, alpha=0.5, c='r', label='Admitted')
16 plt.xlabel('Exam grade 1')
17 plt.ylabel('Exam grade 2')
18
19 x1_min, x1_max = X[:,1].min(), X[:,1].max() # grade 1
20 x2_min, x2_max = X[:,2].min(), X[:,2].max() # grade 2
21
22 xx1, xx2 = np.meshgrid(np.linspace(x1_min, x1_max), np.linspace(x2_min, x2_max)) # create meshgrid
23
24 X2 = np.ones([np.prod(xx1.shape),3])
25 X2[:,1] = xx1.reshape(-1)
26 X2[:,2] = xx2.reshape(-1)
27
28 p = f_pred(X2, w)
29 p = p.reshape(50,50)
30 p2 = f_pred(X2, w_sklearn)
31 p2 = p2.reshape(50,50)
32 plt.contour(xx1, xx2, p, 0, colors='black')
33 plt.contour(xx1, xx2, p2, 0, colors='magenta')
34
35 plt.title('Decision boundary (black with gradient descent and magenta with scikit-learn)')
36 plt.legend()
37 plt.show()
38

```

```

➤ /usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py:760: DataConversionWarning: A column-vector
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:33: UserWarning: No contour levels were found wit
Decision boundary (black with gradient descent and magenta with scikit-learn)

```



## 10. Plot the probability map

```


1 num_a = 110
2 grid_x1 = np.linspace(20,110,num_a)
3 grid_x2 = np.linspace(20,110,num_a)
4
5 score_x1, score_x2 = np.meshgrid( )
6
7 Z = np.zeros( )
8
9 for i in range(len(score_x1)):
10     for j in range(len(score_x2)):
11

```

```

11
12     predict_prob = sigmoid( )
13
14     Z[j, i] = predict_prob
15
16     # actual plotting example
17 fig = plt.figure(figsize=(10,10))
18
19 ax = fig.add_subplot(111)
20 ax.tick_params( )
21 ax.set_xlabel('Exam grade 1')
22 ax.set_ylabel('Exam grade 2')
23
24 ax.set_xlim(20, 110)
25 ax.set_ylim(20, 110)
26
27 cf = ax.contourf( )
28 ax.scatter( )
29 ax.scatter( )
30 cbar = fig.colorbar(cf)
31 cbar.update_ticks()
32
33 plt.legend()
34 plt.show()

```



## Output results

### 1. Plot the dataset in 2D cartesian coordinate system (1pt)

```

1 plt.figure(figsize=(8,8))
2 plt.scatter(x1_idx0, x2_idx0, alpha=0.5, c='b', label='Rejected')
3 plt.scatter(x1_idx1, x2_idx1, alpha=0.5, c='r', label='Admitted')


```



```

3 plt.scatter(x1_ax1, x2_ax1, alpha=0.5, c='r', label='Admitted')
4 plt.title('Training data')
5 plt.xlabel('Exam grade 1')
6 plt.ylabel('Exam grade 2')
7 plt.legend()
8 plt.show()

```




## 2. Plot the sigmoid function (1pt)

```

1 x_values = np.linspace(-10,10)
2
3 plt.figure(2)
4 plt.plot(x_values, sigmoid(x_values))
5 plt.title("Sigmoid function")
6 plt.grid(True)

```




## 3. Plot the loss curve in the course of gradient descent using the mean square error (2pt)

```

1 # plot
2 plt.figure(3)
3 plt.plot([op for op in range(max_iter)], L_iters, c='blue') # plot the loss curve
4 plt.xlabel('Iterations')
5 plt.ylabel('Loss value')


```

```
6 plt.show()
```



▼ 4. Plot the loss curve in the course of gradient descent using the cross-entropy error (2pt)


```
1 # plot
2 plt.figure(3)
3 plt.plot([op for op in range(max_iter)], L_iters, c='blue') # plot the loss curve
4 plt.xlabel('Iterations')
5 plt.ylabel('Loss value')
6 plt.show()
```



▼ 5. Plot the decision boundary using the mean square error (2pt)


```
1 # plot
2 plt.figure(4,figsize=(8,8))
3 plt.scatter(x1_idx0, x2_idx0, alpha=0.5, c='b', label='Rejected')
4 plt.scatter(x1_idx1, x2_idx1, alpha=0.5, c='r', label='Admitted')
5 plt.contour(xx1, xx2, p, 0)
6 plt.xlabel('Exam grade 1')
7 plt.ylabel('Exam grade 2')
8 plt.legend()
9 plt.title('Decision boundary')
10 plt.show()
```





6. Plot the decision boundary using the cross-entropy error (2pt)

```
1 # plot
2 plt.figure(4,figsize=(8,8))
3 plt.scatter(x1_idx0, x2_idx0, alpha=0.5, c='b', label='Rejected')
4 plt.scatter(x1_idx1, x2_idx1, alpha=0.5, c='r', label='Admitted')
5 plt.contour(xx1, xx2, p, 0)
6 plt.xlabel('Exam grade 1')
7 plt.ylabel('Exam grade 2')
8 plt.legend()
9 plt.title('Decision boundary')
10 plt.show()
```



7. Plot the decision boundary using the Scikit-learn logistic regression algorithm (2pt)

1

- ▼ 8. Plot the probability map using the mean square error (2pt)

1

- ▼ 9. Plot the probability map using the cross-entropy error (2pt)

1